

On the ERM Principle with Networked Data

Abstract

Networked data, in which every training example involves two objects and may share some common objects with others, is used in many machine learning tasks such as learning to rank and link prediction. A challenge of learning from networked examples is that target values are not known for some pairs of objects. In this case, neither the classical i.i.d. assumption nor techniques based on complete U -statistics can be used. Most existing theoretical results of this problem only deal with the classical empirical risk minimization (ERM) principle that always weights every example equally, but this strategy leads to unsatisfactory bounds. We consider general weighted ERM and show new universal risk bounds for this problem. These new bounds naturally define an optimization problem which leads to appropriate weights for networked examples. Though this optimization problem is not convex in general, we devise a new fully polynomial-time approximation scheme (FPTAS) to solve it.

1 Introduction

“No man is an island, entire of itself...”, the beginning of a well-known poem by the 17th century English poet John Donne, might be able to explain why social networking websites are so popular. These social media not only make communications convenient and enrich our lives but also bring us data, of an unimaginable amount, that is intrinsically networked. Social network data nowadays is widely used in research on social science, network dynamics, and as an inevitable fate, data mining and machine learning (Scott 2017). Similar examples of networked data such as traffic networks (Min and Wynter 2011), chemical interaction networks (Szkłarczyk et al. 2014), citation networks (Dawson et al. 2014) also abound throughout the machine learning world.

Admittedly, many efforts have been made to design practical algorithms for learning from networked data, e.g., (Liben-Nowell and Kleinberg 2007, Macskassy and Provost 2007, Li et al. 2016, Garcia-Duran et al. 2016). However, not many theoretical guarantees of these methods have been established, which is the main concern of this paper. More specifically, this paper deals with risk bounds of *classifiers trained with networked data* (CLANET) whose goal is to train a classifier with examples in a data graph G . Every vertex of G is an

object and described by a feature vector $X \in \mathcal{X}$ that is drawn independently and identically (i.i.d.) from an unknown distribution, while every edge corresponds to a training example whose input is a pair of feature vectors (X, X') of the two ends of this edge and whose target value Y is in $\{0, 1\}$.

A widely used principle to select a proper model from a hypothesis set is *Empirical Risk Minimization (ERM)*. Papa, Bellet, and Cléménçon (2016) establish risk bounds for ERM on complete data graphs, and the bounds are independent of the distribution of the data. These bounds are of the order $O(\log(n)/n)$, where n is the number of vertices in the complete graph. However, in practice it is very likely that one cannot collect examples for all pairs of vertices and then G is usually incomplete, thus techniques based on complete U -processes in (Papa, Bellet, and Cléménçon 2016) cannot be applied and the risk bounds of the order $O(\log(n)/n)$ are no longer valid in this setting. By generalizing the moment inequality for U -processes to the case of incomplete graphs, we prove novel risk bounds for the incomplete graph.

Usually, every training example is equally weighted (or unweighted) in ERM, which seems much less persuasive when the examples are networked, in particular when the graph is incomplete. But, most existing theoretical results of learning from networked examples are based on the unweighted ERM (Usunier, Amini, and Gallinari 2006, Ralaivola, Szafranski, and Stempfel 2009), and their bounds are of the order $O(\sqrt{\chi^*(D_G)/m})$ where D_G is the *line graph* of G and χ^* is the *fractional chromatic number* of D_G (see Section A in the appendix) and m is the number of training examples. In order to improve this bound, Wang, Guo, and Ramon (2017) propose *weighted ERM* which adds weights to training examples according to the data graph, and show that the risk bound for weighted ERM can be of the order $O(1/\sqrt{\nu^*(G)})$ where $\nu^*(G)$ is the *fractional matching number* of G , so using weighted ERM networked data can be more effectively exploited than the equal weighting method, as basic graph theory tells us $\nu_G^* \geq m/\chi^*(D_G)$. However, Wang, Guo, and Ramon (2017) (in fact, Usunier, Amini, and Gallinari (2006) and Ralaivola, Szafranski, and Stempfel (2009) also) assume that any two examples can be arbitrarily correlated if they share a vertex, which cannot lead to an $O(\log(n)/n)$ bound when the graph is complete.

We show that the “*low-noise*” condition, also called the *Mammen-Tsybakov noise condition* (Mammen and Tsybakov

1998), which is commonly assumed in many typical learning problems with networked data, e.g., ranking (Cl  men  on, Lugosi, and Vayatis 2008) and graph reconstruction (Papa, Bellet, and Cl  men  on 2016), can be used to reasonably bound the dependencies of two examples that share a common vertex and then leads to *tighter* risk bounds.

In summary, in this paper we mainly

- prove new universal risk bounds for CLANET which
 - can be applied to learning from networked data even if the data graph is incomplete;
 - exploit the property of the “low-noise” condition, and then become tighter than previous results;
 - allow non-identical weights on different examples, so it is possible to achieve better learning guarantee by choosing these weights.
- formulate a non-convex optimization problem inspired by our new risk bounds (because our risk bounds depend on the weights added to every training example, and a better weighting scheme leads to a tighter bound), and then we also design a new efficient algorithm to obtain an approximate optimal weighting vector and show that this algorithm is a fully polynomial-time approximation scheme for this non-convex program.

2 Intuitions

We now have a look at previous works that are closely related to our work, as shown in Table 1, and present the merits of our method. Biau and Bleakley (2006), Cl  men  on, Lugosi, and Vayatis (2008) and Papa, Bellet, and Cl  men  on (2016) deal with the case when the graph is complete, i.e., the target value of every pair of vertices is known. In this case, Cl  men  on, Lugosi, and Vayatis (2008) formulate the “low-noise” condition for the ranking problem and demonstrate that this condition can lead to tighter risk bounds by the moment inequality for U -processes. Papa, Bellet, and Cl  men  on (2016) further consider the graph reconstruction problem introduced by Biau and Bleakley (2006) and show this problem always satisfies the “low-noise” condition.

If the graph is incomplete, one can use either Janson’s decomposition (Janson 2004, Usunier, Amini, and Gallinari 2006, Ralaivola, Szafranski, and Stempfel 2009, Ralaivola and Amini 2015) or the fractional matching approach by Wang, Guo, and Ramon (2017) to derive risk bounds. The main differences between these two approaches are:

- Wang, Guo, and Ramon (2017) consider the data graph G while Janson’s decomposition uses only the line graph D_G .
- The fractional matching approach considers weighted ERM while Janson (2004), Usunier, Amini, and Gallinari (2006), Ralaivola, Szafranski, and Stempfel (2009) and Ralaivola and Amini (2015) only prove bounds for unweighted ERM.

Though Wang, Guo, and Ramon (2017) show improved risk bounds, as far as we know, there is no known tight risk

bound on incomplete graphs for tasks such as pairwise ranking and graph reconstruction that satisfy the “low-noise” condition. Under this condition, the method proposed in (Wang, Guo, and Ramon 2017) does not work (see Section 6.1).

Before we show new risk bounds and new weighting methods, we present the following three aspects to convey some intuitions.

Line Graphs Compared to Janson’s decomposition which is based on line graphs, our method utilizes the additional dependency information in the data graph G . For example, the complete line graph with three vertices (i.e., triangle) corresponds to two different data graphs, as illuminated in Figure 1. Hence, line graph based methods ignore some important information in the data graph. This negligence makes it unable to improve bounds, no matter whether considering weighted ERM or not (see Section A.1 in the appendix). In Section 6.2, we show that our bounds are tighter than that of line graph based methods.

Asymptotic Risk As mentioned by Wang, Guo, and Ramon (2017), if several examples share a vertex, then we are likely to put less weight on them because the influence of this vertex to the empirical risk should be bounded. Otherwise, if we treat every example equally, then these dependent examples may dominate the training process and lead to the risk bounds that do not converge to 0 (see the example in Section 6.2).

Uniform Bounds Ralaivola and Amini (2015) prove an entropy-base concentration inequality for networked data using Janson’s decomposition, but the assumption there is usually too restrictive to be satisfied (see Section A.2 in appendix). To circumvent this problem, our method uses the “low-noise” condition (also used in (Papa, Bellet, and Cl  men  on 2016)) to establish uniform bounds, in absence of any restrictive condition imposed on the data distribution.

3 Preliminaries

In this section, we begin with the detailed probabilistic framework for CLANET, and then give the definition of weighted ERM on networked examples.

3.1 Problem Statement

Consider a graph $G = (V, E)$ with a vertex set $V = \{1, \dots, n\}$ and a set of edges $E \subseteq \{\{i, j\} : 1 \leq i \neq j \leq n\}$. For each $i \in V$, a continuous random variable (r.v.) X_i , taking its values in a measurable space \mathcal{X} , describes features of vertex i . The X_i ’s are i.i.d. r.v.’s following some unknown distribution $P_{\mathcal{X}}$. Each pair of vertices $(i, j) \in E$ corresponds to a networked example whose *input* is a pair (X_i, X_j) and *target* value is $Y_{i,j} \in \mathcal{Y}$. We focus on *binary classification* in this paper, i.e., $\mathcal{Y} = \{0, 1\}$. Moreover, the distribution of target values only depends on the features of the vertices it contains but does not depend on features of other vertices, that is, there is a probability distribution $P_{\mathcal{Y}|\mathcal{X}^2}$ such that for

Table 1: Summary of methods for CLANET.

Principles	Graph type	With “low-noise” condition	Without “low-noise” condition
Unweighted ERM (equally weighted)	Complete graphs	Cl��men��on, Lugosi, and Vayatis (Ann. Stat. 2008), Papa, Bellet, and Cl��men��on (NIPS 2016)	Biau and Bleakley (Statistics and Decisions 2006)
	General graphs	Ralaivola and Amini (ICML 2015)	Usunier, Amini, and Gallinari (NIPS 2006), Ralaivola, Szafranski, and Stempfel (AISTATS 2009)
Weighted ERM	General graphs	<i>This paper</i>	Wang, Guo, and Ramon (ALT 2017)

every pair $(i, j) \in E$, the conditional probability

$$P[Y_{i,j} = y \mid x_1, \dots, x_n] = P_{\mathcal{Y}|\mathcal{X}^2}[y, x_i, x_j].$$

Example 1 (pairwise ranking). (Liu 2009) categorize ranking problems into three groups by their input representations and loss functions. One of these categories is pairwise ranking that learns a binary classifier telling which document is better in a given pair of documents. A document can be described by a feature vector from the \mathcal{X} describing title, volume, . . . The target value (rank) between two documents, that only depends on features of these two documents, is 1 if the first document is considered better than the second, and 0 otherwise.

The training set $\mathbb{S} := \{(X_i, X_j, Y_{i,j})\}_{(i,j) \in E}$ is dependent copies of a generic random vector $(X_1, X_2, Y_{1,2})$ whose distribution $P = P_{\mathcal{X}} \otimes P_{\mathcal{X}} \otimes P_{\mathcal{Y}|\mathcal{X}^2}$ is fully determined by the pair $(P_{\mathcal{X}}, P_{\mathcal{Y}|\mathcal{X}^2})$. Let \mathcal{R} be the set of all measurable functions from \mathcal{X}^2 to \mathcal{Y} and for all $r \in \mathcal{R}$, the loss function $\ell(r, (x_1, x_2, y_{1,2})) = \mathbb{1}_{y_{1,2} \neq r(x_1, x_2)}$.

Given a graph G with training examples \mathbb{S} and a hypothesis set $R \subseteq \mathcal{R}$, the CLANET problem is to find a function $r \in R$, with risk

$$L(r) := \mathbb{E}[\ell(r, (X_1, X_2, Y_{1,2}))] \quad (1)$$

that achieves a comparable performance to the Bayes rule $r^* = \arg \inf_{r \in \mathcal{R}} L(r) = \mathbb{1}_{\eta(x_1, x_2) \geq 1/2}$, whose risk is denoted by L^* , where $\eta(x_1, x_2) = P_{\mathcal{Y}|\mathcal{X}^2}[1, x_1, x_2]$ is the regression function.

The main purpose of this paper is to devise a principle to select a classifier \hat{r} from the hypothesis set R and establish bounds for its excess risk $L(\hat{r}) - L^*$.

Definition 1 (“low-noise” condition). Let us consider a learning problem, in which the hypothesis set is \mathcal{F} and the Bayes rule is f^* . With slightly abusing the notation, this problem satisfies the “low-noise” condition if $\forall f \in \mathcal{F}, L(f) - L^* \geq C^\theta (\mathbb{E}[|f - f^*|])^\theta$ where C is a positive constant.

As mentioned, the “low-noise” condition can lead to tighter risk bounds. For this problem, we show that the “low-noise” condition for the i.i.d. part of the Hoeffding decomposition (Hoeffding 1948) of its excess risk can be always obtained if the problem is symmetric (see Lemma 2).

Definition 2 (symmetry). A learning problem is symmetric if for every $x_i, x_j \in \mathcal{X}$, $y_{i,j} \in \mathcal{Y}$ and $r \in \mathcal{R}$, $\ell(r, (x_i, x_j, y_{i,j})) = \ell(r, (x_j, x_i, y_{j,i}))$.

Many typical learning problems are symmetric. For example, pairwise ranking problem with symmetric functions r in the sense that $r(X_1, X_2) = 1 - r(X_2, X_1)$ satisfies the symmetric condition.

3.2 Weighted ERM

ERM aims to find the function from a hypothesis set that minimizes the empirical estimator of (1) on the training examples $\mathbb{S} = \{(X_i, X_j, Y_{i,j})\}_{(i,j) \in E}$:

$$L_m(r) := \frac{1}{m} \sum_{(i,j) \in E} \ell(r, (X_i, X_j, Y_{i,j})). \quad (2)$$

where m is the number of training examples. In this paper, we consider its weighted version, in which we put weights on the examples and select the minimizer $r_{\mathbf{w}}$ of the weighted empirical risk

$$L_{\mathbf{w}}(r) := \frac{1}{\|\mathbf{w}\|_1} \sum_{(i,j) \in E} w_{i,j} \ell(r, (X_i, X_j, Y_{i,j})) \quad (3)$$

where \mathbf{w} is a fractional matching of G and $\|\mathbf{w}\|_1 > 0$.

Definition 3 (fractional matching). Given a graph $G = (V, E)$, a fractional matching \mathbf{w} is a non-negative vector $(w_{i,j})_{(i,j) \in E}$ that for every vertex $i \in V$, $\sum_{j: (i,j) \in E} w_{i,j} \leq 1$.

4 Universal Risk Bounds

In this section, we use covering numbers as the complexity measurement of hypothesis sets to prove that tighter universal risk bounds are always attained by the minimizers of the weighted empirical risk (3).

4.1 Covering Numbers

The excess risk $L(r_{\mathbf{w}}) - L^*$ depends on the hypothesis set R whose complexity can be measured by covering number (Cucker and Zhou 2007). A similar but looser result using VC-dimension (Vapnik and Chervonenkis 1971) can be obtained as well.

Definition 4 (covering numbers). Let $(\mathcal{F}, \mathbb{L}_p)$ be a metric space with \mathbb{L}_p -pseudometric. We define the covering number $N(\mathcal{F}, \mathbb{L}_p, \epsilon)$ be the minimal $l \in \mathbb{N}$ such that there exist l disks in \mathcal{F} with radius ϵ covering \mathcal{F} . If the context is clear, we simply denote $N(\mathcal{F}, \mathbb{L}_p, \epsilon)$ by $N_p(\mathcal{F}, \epsilon)$.

In this paper, we focus on the \mathbb{L}_∞ covering number $N_\infty(\mathcal{F}, \epsilon)$ and suppose that it satisfies the following assumption.

Assumption 1. There exists a nonnegative number $\beta < 1$ and a constant K such that $\log N_\infty(\mathcal{F}, \epsilon) \leq K\epsilon^{-\beta}$ for all $\epsilon \in (0, 1]$.

Similar to (Massart and Nédélec 2006) and (Rejchel 2012), we restrict to $\beta < 1$, whereas in the empirical process theory this exponent usually belongs to $[0, 2)$. This restriction is needed to prove Lemma 1, which involves the integral of $\log N_\infty(\mathcal{F}, \epsilon)$ through 0. Dudley (1974), Korostelev and Tsybakov (1993) and Mammen and Tsybakov (1995) presented various examples of classes \mathcal{F} satisfying Assumption 1. We also refer interested readers to (Mammen and Tsybakov 1998, p. 1813) for more concrete examples of hypothesis classes with smooth boundaries satisfying Assumption 1.

4.2 Risk Bounds

Now we are ready to show the tighter risk bounds for weighted empirical risk by the following theorem.

Theorem 1 (risk bounds). Let $r_{\mathbf{w}}$ be a minimizer of the weighted empirical risk $L_{\mathbf{w}}$ over a class R that satisfies Assumption 1. There exists a constant $C > 0$ such that for all $\delta \in (0, 1]$, with probability at least $1 - \delta$, the excess risk of $r_{\mathbf{w}}$ satisfies

$$L(r_{\mathbf{w}}) - L^* \leq 2\left(\inf_{r \in R} L(r) - L^*\right) + \frac{K'C \log(1/\delta)}{(1 - \beta)^{2/(\beta+1)} \|\mathbf{w}\|_1} \left(\|\mathbf{w}\|_1^{\beta/(1+\beta)} + \max \left(\|\mathbf{w}\|_2, \|\mathbf{w}\|_{\max} (\log(1/\delta))^{1/2}, \|\mathbf{w}\|_{\infty} (\log(1/\delta)) \right) \right) \quad (4)$$

where $\|\mathbf{w}\|_{\max} = \max_i \sqrt{\sum_{j:(i,j) \in E} w_{i,j}^2}$ and $K' = \max(K, \sqrt{K}, K^{1/(1+\beta)})$.

From Theorem 1, for all $\delta \in \left(\exp \left(-\min(\|\mathbf{w}\|_2/\|\mathbf{w}\|_{\infty}, \|\mathbf{w}\|_2^2/\|\mathbf{w}\|_{\max}^2) \right), 1 \right]$, the risk bounds

are of the order $O((1/\|\mathbf{w}\|_1)^{1/(1+\beta)} + \|\mathbf{w}\|_2/\|\mathbf{w}\|_1)$. In this case, our bounds are tighter than $O(1/\sqrt{\|\mathbf{w}\|_1})$ as $\|\mathbf{w}\|_2/\|\mathbf{w}\|_1 \leq 1/\sqrt{\|\mathbf{w}\|_1}$ (recall that \mathbf{w} must be a fractional matching and $0 < \beta < 1$). If G is complete and every example is equally weighted, the bounds of the order $O((1/n)^{1/(1+\beta)})$ achieves the same results in (Papa, Bellet, and Cléménçon 2016) with the same condition on δ but different complexity measurements of the hypothesis set.

Remark. Theorem 1 provides universal risk bounds no matter what the distribution of the data is. The factor of 2 in front of the approximation error $\inf_{r \in R} L(r) - L^*$ has no special meaning and can be replaced by any constant larger than 1 with a cost of increasing the constant C . Wang, Guo, and Ramon (2017) obtain risk bounds that has a factor 1 in front of the approximation error part, but in their result the bound is $O(1/\sqrt{\|\mathbf{w}\|_1})$. Hence, Theorem 1 improves their results if the approximation error does not dominate the other terms in the bounds.

In the rest of this section, we outline the main ideas to obtain this result. We first define

$$q_r(x_1, x_2, y_{1,2}) := \ell(r, x_1, x_2, y_{1,2}) - \ell(r^*, x_1, x_2, y_{1,2})$$

for every $(x_1, x_2, y_{1,2}) \in \mathcal{X} \times \mathcal{X} \times \mathcal{Y}$ and let $\Lambda(r) := L(r) - L^* = \mathbb{E}[q_r(X_1, X_2, Y_{1,2})]$ be the excess risk with respect to the Bayes rule. Its empirical estimate by weighted ERM is

$$\begin{aligned} \Lambda_{\mathbf{w}}(r) &= L_{\mathbf{w}}(r) - L_{\mathbf{w}}(r^*) \\ &= \frac{1}{\|\mathbf{w}\|_1} \sum_{(i,j) \in E} w_{i,j} q_r(X_i, X_j, Y_{i,j}). \end{aligned}$$

By Hoeffding's decomposition (Hoeffding 1948), for all $r \in \mathcal{R}$, one can write

$$\Lambda_{\mathbf{w}}(r) = T_{\mathbf{w}}(r) + U_{\mathbf{w}}(r) + \tilde{U}_{\mathbf{w}}(r), \quad (5)$$

where

$$T_{\mathbf{w}}(r) = \Lambda(r) + \frac{2}{\|\mathbf{w}\|_1} \sum_{i=1}^n \sum_{j:(i,j) \in E} w_{i,j} h_r(X_i)$$

is a weighted average of i.i.d. random variables with $h_r(X_i) = \mathbb{E}[q_r(X_i, X_j, Y_{i,j}) \mid X_i] - \Lambda(r)$,

$$U_{\mathbf{w}}(r) = \frac{1}{\|\mathbf{w}\|_1} \sum_{(i,j) \in E} w_{i,j} (\hat{h}_r(X_i, X_j) - h_r(X_i) - h_r(X_j))$$

is a weighted *degenerated* (i.e., the symmetric kernel $\hat{h}_r(x_1, x_2)$ such that $\mathbb{E}[\hat{h}_r(X_1, X_2) \mid X_1 = x_1] = 0$ for all $x_1 \in \mathcal{X}$) U -statistic $\hat{h}_r(X_i, X_j) = \mathbb{E}[q_r(X_i, X_j, Y_{i,j}) \mid X_i, X_j] - \Lambda(r) - h_r(X_i) - h_r(X_j)$ and

$$\tilde{U}_{\mathbf{w}}(r) = \frac{1}{\|\mathbf{w}\|_1} \sum_{(i,j) \in E} w_{i,j} \tilde{h}_r(X_i, X_j, Y_{i,j})$$

with a *degenerated* kernel $\tilde{h}_r(X_i, X_j, Y_{i,j}) = q_r(X_i, X_j, Y_{i,j}) - \mathbb{E}[q_r(X_i, X_j, Y_{i,j}) \mid X_i, X_j]$. In the following, we bound the three terms $T_{\mathbf{w}}$, $U_{\mathbf{w}}$ and $\tilde{U}_{\mathbf{w}}$ in (5) respectively.

Lemma 1 (uniform approximation). Under the same assumptions as in Theorem 1, for any $\delta \in (0, 1/e)$, we have with probability at least $1 - \delta$,

$$\begin{aligned} \sup_{r \in R} |U_{\mathbf{w}}(r)| &\leq \frac{\max(K, \sqrt{K}) C_1}{1 - \beta} \max \left(\frac{\|\mathbf{w}\|_2 \log(1/\delta)}{\|\mathbf{w}\|_1}, \right. \\ &\quad \left. \frac{\|\mathbf{w}\|_{\max} (\log(1/\delta))^{3/2}}{\|\mathbf{w}\|_1}, \frac{\|\mathbf{w}\|_{\infty} (\log(1/\delta))^2}{\|\mathbf{w}\|_1} \right) \end{aligned}$$

and

$$\sup_{r \in R} |\tilde{U}_{\mathbf{w}}(r)| \leq \frac{\max(K, \sqrt{K})C_2}{1 - \beta} \left(\frac{\|\mathbf{w}\|_2}{\|\mathbf{w}\|_1} + \max \left(\frac{\|\mathbf{w}\|_{\max}(\log(1/\delta))^{3/2}}{\|\mathbf{w}\|_1}, \frac{\|\mathbf{w}\|_{\infty}(\log(1/\delta))^2}{\|\mathbf{w}\|_1} \right) \right)$$

where $C_1, C_2 < +\infty$ are constants.

To prove Lemma 1, we show that $U_{\mathbf{w}}(r)$ and $\tilde{U}_{\mathbf{w}}(r)$ can be bounded by Rademacher chaos using classical symmetrization and randomization tricks combined with the decoupling method. We handle these Rademacher chaos by generalizing the moment inequality for U -statistics in (Cl  men  on, Lugosi, and Vayatis 2008). Specifically, we utilize the moment inequalities from (Boucheron et al. 2005) to convert them into a sum of simpler processes, which can be bounded by the metric entropy inequality for Khinchine-type processes (see Arcones and Gine 1993, Proposition 2.6) and Assumption 1. The detailed proofs can be found in Section C in the appendix.

Lemma 1 shows that the contribution of the degenerated parts $U_{\mathbf{w}}(r)$ and $\tilde{U}_{\mathbf{w}}(r)$ to the excess risk can be bounded. This implies that minimizing $\Lambda_{\mathbf{w}}(r)$ is approximately equivalent to minimizing $T_{\mathbf{w}}(r)$ and thus $r_{\mathbf{w}}$ is a ρ -minimizer of $T_{\mathbf{w}}(r)$ in the sense that $T_{\mathbf{w}}(r_{\mathbf{w}}) \leq \rho + \inf_{r \in R} T_{\mathbf{w}}(r)$. In order to analyze $T_{\mathbf{w}}(r)$, which can be treated as a weighted empirical risk on i.i.d. examples, we generalize the results in (Massart and N  d  lec 2006) (see Section B in the appendix). Based on this result, tight bounds for the excess risk with respect to $T_{\mathbf{w}}(r)$ can be obtained if the variance of the excess risk is controlled by its expected value. By Lemma 2, $T_{\mathbf{w}}(r)$ fulfills this condition, which leads to Lemma 3.

Lemma 2 (condition leads to ‘‘low-noise’’, (Papa, Bellet, and Cl  men  on 2016, Lemma 2)). *If the learning problem CLANET is symmetric, then*

$$\text{Var}[\mathbb{E}[q_r(X_1, X_2, Y_{1,2}) \mid X_1]] \leq \Lambda(r) \quad (6)$$

holds for any distribution P and any function $r \in R$.

Lemma 3 (risk bounds for i.i.d. examples). *Suppose that r' is a ρ -minimizer of $T_{\mathbf{w}}(r)$ in the sense that $T_{\mathbf{w}}(r') \leq \rho + \inf_{r \in R} T_{\mathbf{w}}(r)$ and R satisfies Assumption 1, then there exists a constant C such that for all $\delta \in (0, 1]$, with probability at least $1 - \delta$, the risk of r' satisfies*

$$\Lambda(r') \leq 2 \inf_{r \in R} \Lambda(r) + 2\rho + \frac{CK^{1/(1+\beta)} \log(1/\delta)}{(\|\mathbf{w}\|_1(1 - \beta)^2)^{1/(1+\beta)}}.$$

With Lemma 1 Lemma 3, now we are ready to prove Theorem 1.

Proof of Theorem 1. Let us consider the Hoeffding decomposition (5) of $\Lambda_{\mathbf{w}}(r)$ that is minimized over $r \in R$. The idea of this proof is that the degenerate parts $U_{\mathbf{w}}(r)$ and $\tilde{U}_{\mathbf{w}}(r)$ can be bounded by Lemma 1. Therefore, $r_{\mathbf{w}}$ is an approximate minimizer of $T_{\mathbf{w}}(r)$, which can be handled by Lemma 3.

Let A be the event that

$$\sup_{r \in R} |U_{\mathbf{w}}(r)| \leq \kappa_1,$$

where

$$\kappa_1 = \frac{C_1}{1 - \beta} \max \left(\frac{\|\mathbf{w}\|_2 \log(1/\delta)}{\|\mathbf{w}\|_1}, \frac{\|\mathbf{w}\|_{\max}(\log(1/\delta))^{3/2}}{\|\mathbf{w}\|_1}, \frac{\|\mathbf{w}\|_{\infty}(\log(1/\delta))^2}{\|\mathbf{w}\|_1} \right)$$

for an appropriate constant C_1 . Then by Lemma 1, $P[A] \geq 1 - \delta/4$. Similarly, let B be the event that

$$\sup_{r \in R} |\tilde{U}_{\mathbf{w}}(r)| \leq \kappa_2.$$

where

$$\kappa_2 = \frac{C_2}{1 - \beta} \left(\frac{\|\mathbf{w}\|_2}{\|\mathbf{w}\|_1} + \max \left(\frac{\|\mathbf{w}\|_{\max}(\log(1/\delta))^{3/2}}{\|\mathbf{w}\|_1}, \frac{\|\mathbf{w}\|_{\infty}(\log(1/\delta))^2}{\|\mathbf{w}\|_1} \right) \right)$$

for an appropriate constant C_2 . Then $P[B] \geq 1 - \delta/4$.

By (5), it is clear that, if both A and B happen, $r_{\mathbf{w}}$ is a ρ -minimizer of $T_{\mathbf{w}}(r)$ over $r \in R$ in the sense that the difference between the value of this latter quantity at its minimum and $r_{\mathbf{w}}$ is at most $(\kappa_1 + \kappa_2)$. Then, from Lemma 3, with probability at least $1 - \delta/2$, $r_{\mathbf{w}}$ is a $(\kappa_1 + \kappa_2)$ -minimizer of $T_{\mathbf{w}}(r)$, which the result follows. \square

An intuition obtained from our result is how to choose weights for networked data. By Theorem 1, to obtain tight risk bounds, we need to maximize $\|\mathbf{w}\|_1$ (under the constraint that this weight vector is a fractional matching), which resembles the result of (Wang, Guo, and Ramon 2017) (but they only need to maximize $\|\mathbf{w}\|_1$ and this is why they end in the $O(1/\sqrt{\nu^*(G)})$ bound), while making $\|\mathbf{w}\|_2, \|\mathbf{w}\|_{\max}, \|\mathbf{w}\|_{\infty}$ as small as possible, which appears to suggest putting nearly average weights on examples and vertices respectively. These two objectives, maximizing $\|\mathbf{w}\|_1$ and minimizing $\|\mathbf{w}\|_2, \|\mathbf{w}\|_{\max}, \|\mathbf{w}\|_{\infty}$, seem to contradict each other. In the next section, we discuss how to solve this problem.

5 Weighting Vector Optimization

In this section, we first formulate the optimization problem that minimizes the risk bounds in Theorem 1. Although this optimization problem is not convex unless $\beta = 0$, which usually means that there is no general efficient way to solve it, we devise a *fully polynomial-time approximation scheme* (FPTAS) to solve it.

Definition 5 (FPTAS). *An algorithm \mathcal{A} is a FPTAS for a minimization problem Π , if for any input \mathcal{I} of Π and $\epsilon > 0$, \mathcal{A} finds a solution s in time polynomial in both the size of \mathcal{I} and $1/\epsilon$ that satisfies*

$$f_{\Pi}(s) \leq (1 + \epsilon) \cdot f_{\Pi}(s^*),$$

where f_{Π} is the (positive) objective function of Π and s^* is an optimal solution for \mathcal{I} .

5.1 Optimization Problem

According to Theorem 1, given a graph G , $\beta \in (0, 1)$ and $\delta \in (0, 1]$, one can find a good weighting vector with tight risk bounds by solving the following program:

$$\begin{aligned} \min_{\mathbf{w}} \quad & \frac{1}{\|\mathbf{w}\|_1} \left(\|\mathbf{w}\|_1^{\beta/(1+\beta)} + \max \left(\|\mathbf{w}\|_2, \right. \right. \\ & \left. \left. \|\mathbf{w}\|_{\max} (\log(1/\delta))^{1/2}, \|\mathbf{w}\|_{\infty} (\log(1/\delta)) \right) \right) \\ \text{s.t.} \quad & \forall (i, j) \in E, w_{i,j} \geq 0 \quad \text{and} \quad \forall i, \sum_{j:(i,j) \in E} w_{i,j} \leq 1 \end{aligned} \quad (7)$$

To get rid of the fraction of norms in the program above, we consider a distribution \mathbf{p} on edges $p_{i,j} := w_{i,j}/\|\mathbf{w}\|_1$ and then $\|\mathbf{w}\|_1 \leq 1/\max_{i=1,\dots,n} \sum_{j:(i,j) \in E} p_{i,j}$. Every distribution \mathbf{p} corresponds to a valid weighting vector \mathbf{w} . By introducing two auxiliary variables a and b , solving the original program (7) is equivalent to solving

$$\begin{aligned} \min_{a,b,\mathbf{p}} \quad & a^{1/(1+\beta)} + b \\ \text{s.t.} \quad & \forall (i, j) \in E, p_{i,j} \geq 0 \\ & \forall (i, j) \in E, p_{i,j} \log(1/\delta) - b \leq 0 \\ & \forall i, \sum_{j:(i,j) \in E} p_{i,j} - a \leq 0 \\ & \forall i, \left(\sum_{j:(i,j) \in E} p_{i,j}^2 \log(1/\delta) \right)^{1/2} - b \leq 0 \\ & \|\mathbf{p}\|_2 - b \leq 0 \quad \text{and} \quad \sum_{(i,j) \in E} p_{i,j} = 1 \end{aligned} \quad (8)$$

Note that the constraints are all convex. If $\beta = 0$, e.g., the hypothesis set is finite, then the objective function becomes linear and thus (8) is a convex optimization problem that can be solved by some convex optimization method (see e.g., (Boyd and Vandenberghe 2004)) such as interior-point method.

If $\beta > 0$, the objective function is not convex any more. In fact, the program (8) becomes a concave problem that may be optimized globally by some complex algorithms (Benson 1995, Hoffman 1981) that often need tremendous computation. Instead, one may only need to approximate it using some efficient methods, e.g., Concave-Convex Procedure (Yuille 2001) and Coordinate Descent (Wright 2015). However, these methods lack in complexity analysis and may lead to a local optimum.

5.2 A Fully Polynomial-time Approximation Scheme

To solve the program (8) efficiently, we propose Algorithm 1 and show that it is a fully polynomial-time approximation scheme for (8).

¹For example, some interior-point method.

Algorithm 1 FPTAS for weighting vector optimization.

Input: ϵ, β, δ and a graph G that contains n vertices and m edges.

Output: An approximate optimal weighting vector $\bar{\mathbf{p}}$ for the program (8).

1: Solve the following linear program (LP) efficiently¹, and obtain an ϵ -approximation a_{\min} ;

$$\begin{aligned} \min_{a,\mathbf{p}} \quad & a \\ \text{s.t.} \quad & \forall (i, j) \in E, p_{i,j} \geq 0 \\ & \forall i, \sum_{j:(i,j) \in E} p_{i,j} - a \leq 0 \\ & \sum_{(i,j) \in E} p_{i,j} = 1 \end{aligned} \quad (9)$$

2: Let $Grid := \{a_{\min} + i \cdot \epsilon(1 + \beta)/n \mid i \in \mathbb{N} \text{ and } i \leq n(1 - a_{\min})/\epsilon(1 + \beta)\}$ and $Solutions := \emptyset$.

3: **for** $a \in Grid$ **do**

4: Use some efficient interior-point method to obtain an ϵ -approximation of the following program and add the solution (a, b, \mathbf{p}) into $Solutions$.

$$\begin{aligned} \min_{b,\mathbf{p}} \quad & b \\ \text{s.t.} \quad & \forall (i, j) \in E, p_{i,j} \geq 0 \\ & \forall (i, j) \in E, p_{i,j} \log(1/\delta) - b \leq 0 \\ & \forall i, \sum_{j:(i,j) \in E} p_{i,j} - a \leq 0 \\ & \forall i, \left(\sum_{j:(i,j) \in E} p_{i,j}^2 \log(1/\delta) \right)^{1/2} - b \leq 0 \\ & \|\mathbf{p}\|_2 - b \leq 0 \quad \text{and} \quad \sum_{(i,j) \in E} p_{i,j} = 1 \end{aligned} \quad (10)$$

5: **return** the vector $\bar{\mathbf{p}}$ which makes $a^{1/(1+\beta)} + b$ smallest from $Solutions$.

Theorem 2. *Algorithm 1 is a FPTAS for the program (8).*

Proof. We first analyze the running time of this algorithm.

Note that $1/n \leq a \leq 1$ if the graph is not empty. In Algorithm 1, we first divide the problem into at most $\frac{1-1/n}{\epsilon(1+\beta)/n} = \frac{n-1}{\epsilon(1+\beta)}$ convex programs, each of which produces an ϵ -approximate solution by some interior-point method. Since interior-point method is FPTAS for convex problems (Boyd and Vandenberghe 2004), solving each of these programs needs polynomial time in the problem size $m + n$ and $1/\epsilon$. Thus, the complexity of Algorithm 1 is also polynomial in $m + n$ and $1/\epsilon$.

Now we show that this algorithm indeed results in an ϵ -approximation of this optimal solution.

For any optimal solution (a^*, b^*, \mathbf{p}^*) , if a^* achieves minimum for the program (9), we can find a' in $Grid$ (actually

$a_{\min})$ such that

$$\begin{aligned} (a')^{1/(1+\beta)} &\leq (1+\epsilon)^{1/(1+\beta)} (a^*)^{1/(1+\beta)} \\ &\leq (1+\epsilon) (a^*)^{1/(1+\beta)}. \end{aligned} \quad (11)$$

Otherwise, we can also find a' in *Grid* such that $a^* \leq a' < a^* + \epsilon(1+\beta)/n$ and thus

$$\begin{aligned} (a')^{1/(1+\beta)} &\leq (a^* + \epsilon(1+\beta)/n)^{1/(1+\beta)} \\ &\leq (a^*)^{1/(1+\beta)} + \epsilon(1+\beta)/n \\ &\quad \frac{1}{1+\beta} (a^*)^{-\beta/(1+\beta)} \\ &\leq (1+\epsilon) (a^*)^{1/(1+\beta)} \end{aligned} \quad (12)$$

The third inequality follows from the fact that $1/n \leq a^*$. We assume that the optimal solution for the program (10) is $b = b'$ when we fix $a = a'$. Because (a^*, b^*) is feasible and $a' > a^*$, (a', b^*) is always a feasible solution for the program (10), which leads to $b' \leq b^*$. Besides, interior-point method can produce an ϵ -approximate solution b'' such that

$$b'' \leq (1+\epsilon)b' \leq (1+\epsilon)b^*. \quad (13)$$

Finally, we select the best approximate weighting vector $\bar{\mathbf{p}}$ from all solutions in *Solutions*. Combining (11), (12) and (13), we have the objective value for $\bar{\mathbf{p}}$

$$\begin{aligned} (a_{\bar{\mathbf{p}}})^{1/(1+\beta)} + b_{\bar{\mathbf{p}}} &\leq (a')^{1/(1+\beta)} + b'' \\ &\leq (1+\epsilon) ((a^*)^{1/(1+\beta)} + b^*). \end{aligned}$$

□

6 Discussion

In this section, we first show that, according to our bounds, equal weighting is indeed the best weighting scheme for complete graphs. Then, we discuss the performance of this equal weighting scheme when the graph is incomplete.

6.1 Complete Graphs

When graph G is complete, weighting all examples equally gives the best risk bound, as all the terms $\max_{i=1,\dots,n} \sum_{j:(i,j) \in E} p_{i,j}$, $\|\mathbf{p}\|_2$, $\|\mathbf{p}\|_{\max}$ and $\|\mathbf{p}\|_{\infty}$ achieve minimum. Compared to the results in (Wang, Guo, and Ramon 2017), our theory puts additional constraints on $\|\mathbf{p}\|_2$, $\|\mathbf{p}\|_{\max}$ and $\|\mathbf{p}\|_{\infty}$ which encourages weighting examples fairly in this case, as illustrated in Figure 1. Besides, this scheme, which coincides with U -statistics that *average* the basic estimator applied to all sub-samples, produces the smallest variance among all unbiased estimators (Hoeffding 1948).

6.2 Equal Weighting

Let us discuss further the equal weighting scheme that gives every example the same weight. Denote by $\Delta(G)$ the maximum degree of G (note that this is not the maximum degree of D_G) and let $p_{i,j} = 1/m$ (recall that m is the number of examples) for all $(i,j) \in E$. According to program (8), using equal weighting scheme, the risk bounds are of the order

$$O\left(\left(\frac{\Delta(G)}{m}\right)^{1/(1+\beta)} + \frac{1}{\sqrt{m}}\right), \quad (14)$$




Figure 1: (a) and (b) are two different data graphs, but both of them correspond to the same line graph (a triangle). (c) and (d) are two weighting schemes for a complete graph formed by points X_1, X_2, X_3, X_4 . Solid line means its weight $p > 0$ while dash line means $p = 0$. (c): Weight every example equally. (d): Only the two examples in an independent subset get equally non-zero weights and other weights are 0 (dashed line). Note that $\max_{i=1,\dots,n} \sum_{j:(i,j) \in E} p_{i,j}$ of these two weighting schemes are the same, but (c) has the tighter risk bounds, as $\|\mathbf{p}\|_2$, $\|\mathbf{p}\|_{\max}$ and $\|\mathbf{p}\|_{\infty}$ of (c) are smaller than that of (d) respectively.

if $\delta \in (\exp(-m/\Delta(G)), 1]$. In some cases, such as bounded degree graphs and complete graphs, this scheme provides reasonable risk bounds. Note that $\Delta(G)$ is smaller than the maximum size of cliques in its corresponding line graph D_G and $\chi^*(D_G)$ is larger than the maximum size of cliques in D_G , these bounds above are always better than the bounds of the order $O(\sqrt{\chi^*(D_G)/m})$ built by Janson's decomposition.

However, as argued in Section 2, one can construct examples to illustrate that if we use the equal weighting strategy when $\Delta(G)$ is large (e.g., if it is linear to m), the risk bounds (14) are very large and do not converge to 0, while this problem can be solved by simply using a better weighting strategy.

Example 2. Consider a data graph with $|E| = m \gg 1$ and E consists of $m/2$ disjoint edges and $m/2$ edges sharing a common vertex, then $\Delta(G) = m/2$. Using the equal weighting scheme, the risk bounds are of the order $O(1)$ that is meaningless. A much better weighting scheme of this case is to weight the examples of disjoint edges with $2/(m+2)$ while weight the examples of adjacent edges with $4/(m(m+2))$, which provides risk bounds of the order $O\left((1/m)^{1/(1+\beta)} + \sqrt{1/m}\right)$.

7 Conclusion

In this paper, we consider weighted ERM of the symmetric CLANET problem and establish new universal risk bounds under the “low-noise” condition. These new bounds are tighter in the case of incomplete graphs and can be degenerate to the known tightest bound when graphs are complete. Based on this result, one can train a classifier with a better risk bound by putting proper weights on training examples. We propose an efficient algorithm to obtain the approximate optimal weighting vector and prove that the algorithm is a FPTAS for the weighting vector optimization problem. Finally, we discuss two cases to show the merits of our new risk bounds. In particular, we show that, in terms of these risk bounds, weighting all examples equally is optimal for complete graphs but not always good for incomplete graphs.

References

- Arcones, M. A., and Gine, E. 1993. Limit theorems for u-processes. *The Annals of Probability* 1494–1542.
- Benson, H. P. 1995. *Concave Minimization: Theory, Applications and Algorithms*. Springer US.
- Biau, G., and Bleakley, K. 2006. Statistical inference on graphs. *Statistics & Decisions* 24(2):209–232.
- Boucheron, S.; Bousquet, O.; Lugosi, G.; Massart, P.; et al. 2005. Moment inequalities for functions of independent random variables. *The Annals of Probability* 33(2):514–560.
- Bousquet, O. 2002. A bennett concentration inequality and its application to suprema of empirical processes. *Comptes Rendus Mathématique* 334(6):495–500.
- Boyd, S., and Vandenberghe, L. 2004. *Convex optimization*. Cambridge university press.
- Cléménçon, S.; Lugosi, G.; and Vayatis, N. 2008. Ranking and empirical minimization of u-statistics. *Annals of Statistics* 36(2):844–874.
- Cucker, F., and Zhou, D. X. 2007. *Learning theory: an approximation theory viewpoint*, volume 24. Cambridge University Press.
- Dawson, S.; Gašević, D.; Siemens, G.; and Joksimovic, S. 2014. Current state and future trends: A citation network analysis of the learning analytics field. In *Proceedings of the Fourth International Conference on Learning Analytics and Knowledge*, 231–240.
- De la Pena, V., and Giné, E. 2012. *Decoupling: from dependence to independence*. Springer Science & Business Media.
- Dudley, R. M. 1974. Metric entropy of some classes of sets with differentiable boundaries. *Journal of Approximation Theory* 10(3):227–236.
- Garcia-Duran, A.; Bordes, A.; Usunier, N.; and Grandvalet, Y. 2016. Combining two and three-way embedding models for link prediction in knowledge bases. *Journal of Artificial Intelligence Research* 55:715–742.
- Hoeffding, W. 1948. A class of statistics with asymptotically normal distribution. *The Annals of Mathematical Statistics* 293–325.
- Hoffman, K. L. 1981. A method for globally minimizing concave functions over convex sets. *Mathematical Programming* 20(1):22–32.
- Janson, S. 2004. Large deviations for sums of partly dependent random variables. *Random Structures & Algorithms* 24(3):234–248.
- Korostelev, A. P., and Tsybakov, A. B. 1993. *Minimax Theory of Image Reconstruction*. Springer-Verlag.
- Li, J.; Hu, X.; Wu, L.; and Liu, H. 2016. Robust unsupervised feature selection on networked data. In *Proceedings of the 2016 SIAM International Conference on Data Mining*, 387–395.
- Liben-Nowell, D., and Kleinberg, J. 2007. The link-prediction problem for social networks. *Journal of the Association for Information Science and Technology* 58(7):1019–1031.
- Liu, T.-Y. 2009. Learning to rank for information retrieval. *Foundations and Trends® in Information Retrieval* 3(3):225–331.
- Macskassy, S. A., and Provost, F. 2007. Classification in networked data: A toolkit and a univariate case study. *Journal of Machine Learning Research* 8(May):935–983.
- Mammen, E., and Tsybakov, A. B. 1995. Asymptotical minimax recovery of sets with smooth boundaries. *Annals of Statistics* 23(2):502–524.
- Mammen, E., and Tsybakov, A. B. 1998. Smooth discrimination analysis. *Annals of Statistics* 27(6):1808–1829.
- Massart, P., and Nédélec, É. 2006. Risk bounds for statistical learning. *Annals of Statistics* 2326–2366.
- Min, W., and Wynter, L. 2011. Real-time road traffic prediction with spatio-temporal correlations. *Transportation Research Part C: Emerging Technologies* 19(4):606–616.
- Papa, G.; Bellet, A.; and Cléménçon, S. 2016. On graph reconstruction via empirical risk minimization: Fast learning rates and scalability. In *Advances in Neural Information Processing Systems*, 694–702.
- Ralaivola, L., and Amini, M.-R. 2015. Entropy-based concentration inequalities for dependent variables. In *International Conference on Machine Learning*, 2436–2444.
- Ralaivola, L.; Szafranski, M.; and Stempfel, G. 2009. Chromatic pac-bayes bounds for non-iid data. In *Artificial Intelligence and Statistics*, 416–423.
- Rejchel, W. 2012. On ranking and generalization bounds. *Journal of Machine Learning Research* 13(May):1373–1392.
- Scott, J. 2017. *Social network analysis*. Sage.
- Szklarczyk, D.; Franceschini, A.; Wyder, S.; Forslund, K.; Heller, D.; Huerta-Cepas, J.; Simonovic, M.; Roth, A.; Santos, A.; Tsafou, K. P.; et al. 2014. String v10: protein–protein interaction networks, integrated over the tree of life. *Nucleic Acids Research* 43(D1):D447–D452.
- Tsybakov, A. B. 2004. Optimal aggregation of classifiers in statistical learning. *Annals of Statistics* 135–166.
- Usunier, N.; Amini, M.-R.; and Gallinari, P. 2006. Generalization error bounds for classifiers trained with interdependent data. In *Advances in neural information processing systems*, 1369–1376.
- Vapnik, V., and Chervonenkis, A. Y. 1971. On the uniform convergence of relative frequencies of events to their probabilities. *Theory of Probability & Its Applications* 16(2):264–280.
- Wang, Y.; Guo, Z.-C.; and Ramon, J. 2017. Learning from networked examples. In *Proceedings of the 28th International Conference on Algorithmic Learning Theory*, to appear.
- Wright, S. J. 2015. Coordinate descent algorithms. *Mathematical Programming* 151(1):3–34.
- Yuille, A. L. 2001. The concave-convex procedure (cccp). In *Advances in Neural Information Processing Systems 14*, 915 – 936.